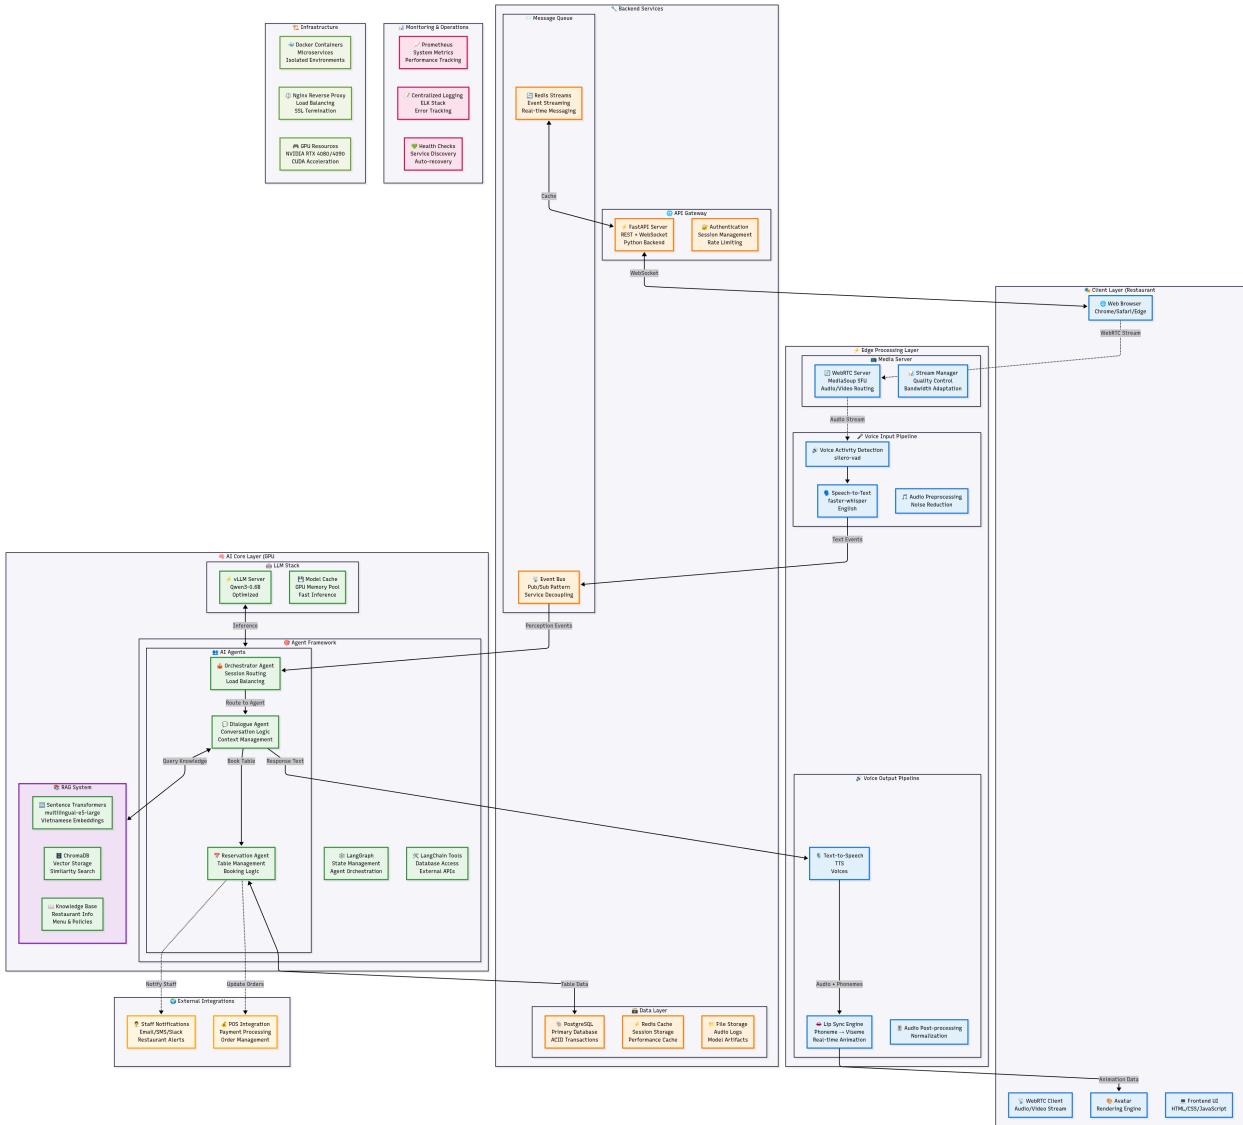




# Tech Stack Architecture

Type	Document
Status	Under Review
Priority	Urgent
Assigned To	Thinh Hung Ho  Minh Hoàng
Due Date	@September 10, 2025
Dependencies	None
Description	Comprehensive documentation of the entire receptionist system architecture, including component interactions, data flows, and integration points.

## Tech Stack Architecture Workflow



# Tech Stack Architecture & Source Structure

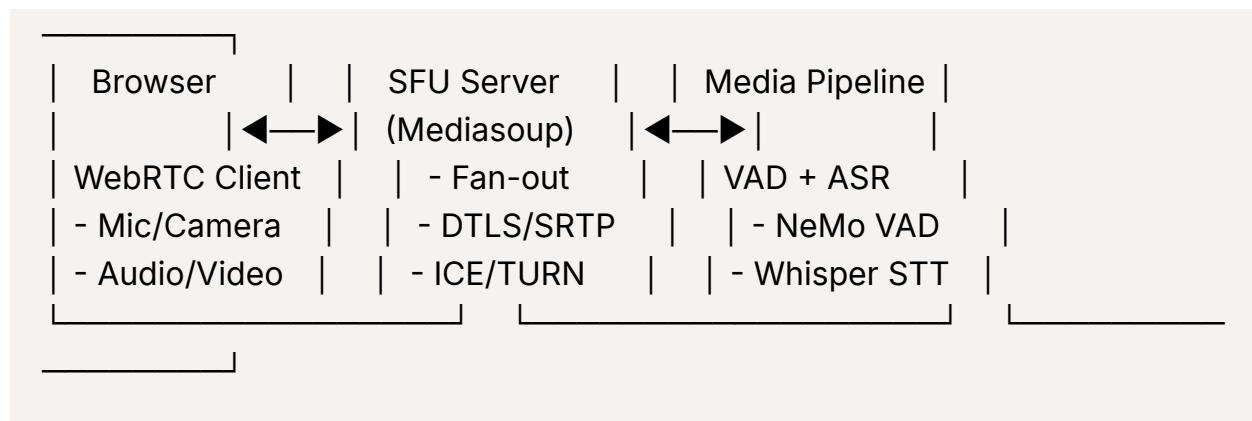
## 1. Detailed Tech Stack Architecture

### 1.1 Perception Pipeline Stack

#### WebRTC & Media Processing

Edge Node Architecture:





## Voice Activity Detection (VAD)

- **Primary:** NeMo MarbleNet (NVIDIA)
- **Fallback:** WebRTC VAD or Silero VAD
- **Output:** Audio segments with confidence scores
- **Latency Target:** <50ms detection

## Automatic Speech Recognition (ASR)

- **Primary:** OpenAI Whisper (streaming adaptation)
- **Alternative:** Google Speech-to-Text Streaming API
- **Features:** Timestamps, confidence scores, Vietnamese + English
- **Latency Target:** <200ms for short utterances

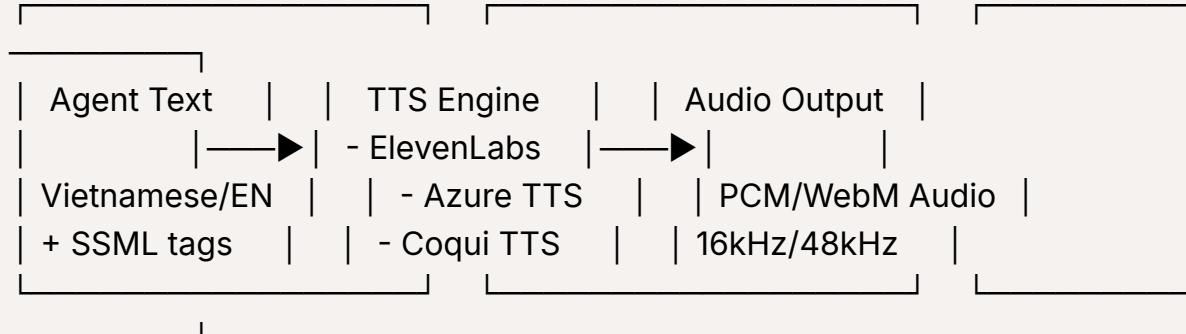
## Speaker Diarization

- **Model:** ECAPA-TDNN embeddings
- **Clustering:** Online clustering (incremental)
- **Output:** Speaker IDs with timestamp ranges

## 1.2 Avatar Pipeline Stack

### Text-to-Speech (TTS)

### TTS Pipeline:



## Lip Sync & Animation

- **3D Avatar:** Ready Player Me or custom Three.js model
- **Lip Sync:** Rhubarb Lip Sync or OVRLipSync
- **Animation:** Mixamo animations + facial expressions
- **Rendering:** Three.js WebGL in browser

## WebRTC Output

- **Codec:** Opus for audio, VP9/H.264 for video
- **Streaming:** MediaStream API to WebRTC peer connection
- **Latency Target:** <150ms end-to-end

## 1.3 Agent Layer Stack (LangGraph + LangChain)

### LangGraph State Machine Architecture

#### # Agent State Flow

##### Orchestrator Agent:

- └── SessionRouter (route events by session\_id)
- └── BackpressureController (rate limiting)
- └── HealthChecker (timeout/retry logic)
- └── EventDispatcher (emit to dialogue agent)

#### Dialogue Agent:

- └─ IntentClassifier (reservation/faq/chitchat)
- └─ ContextManager (session memory)
- └─ ToolSelector (which tools to call)
- └─ ResponseGenerator (craft natural replies)
- └─ GuardrailsChecker (safety/policy validation)

#### Reservation Agent:

- └─ TableMatcher (find available tables)
- └─ HoldManager (temporary reservations)
- └─ AlternativeGenerator (backup options)
- └─ StaffNotifier (alert restaurant staff)
- └─ BookingConfirmmer (finalize reservation)

## 1.4 Backend Services Stack

### Message Bus

- **Primary:** Apache Kafka (for durability)
- **Alternative:** Redis Streams (for simplicity)
- **Topics:** `perception-events` , `agent-utterances` , `booking-actions`

### Database Layer

-- PostgreSQL with pgvector extension

#### Tables:

- └─ sessions (active conversations)
- └─ reservations (booking data)
- └─ tables (restaurant layout/availability)
- └─ knowledge\_base (FAQ embeddings)
- └─ conversation\_logs (audit trail)
- └─ user\_profiles (preferences/history)

### Vector Database (RAG)

- **Primary:** Postgres + pgvector
- **Alternative:** Qdrant or FAISS
- **Embeddings:** OpenAI text-embedding-ada-002
- **Use Cases:** Menu items, policies, common questions

## 2. Source Structure

```

src/
├── README.md
├── docker-compose.yml
├── .env.example
├── .gitignore
└── services/
    ├── edge-node/          # Media processing at edge
    │   ├── Dockerfile
    │   ├── requirements.txt
    │   └── src/
    │       ├── __init__.py
    │       ├── main.py        # FastAPI app entry point
    │       └── webrtc/
    │           ├── __init__.py
    │           ├── sfu_server.py # Mediasoup SFU wrapper
    │           ├── ice_config.py # STUN/TURN configuration
    │           └── peer_manager.py # WebRTC peer connections
    │       └── perception/
    │           ├── __init__.py
    │           ├── pipeline.py  # Main perception orchestrator
    │           └── vad/
    │               ├── __init__.py
    │               ├── nemo_vad.py # NeMo MarbleNet integration
    │               └── fallback_vad.py
    │           └── asr/
    │               ├── __init__.py
    │               └── whisper_streaming.py

```

```
    |   |   |   └── google_stt.py
    |   |   └── diarization/
    |   |       ├── __init__.py
    |   |       ├── ecapa_tdnn.py
    |   |       └── clustering.py
    |   └── emotion/
    |       ├── __init__.py
    |       └── prosody_analyzer.py
    └── avatar/
        ├── __init__.py
        ├── pipeline.py      # Avatar rendering orchestrator
        └── tts/
            ├── __init__.py
            ├── elevenlabs_tts.py
            ├── azure_tts.py
            └── coqui_tts.py
        └── lipsync/
            ├── __init__.py
            ├── rhubarb_sync.py
            └── ovr_lipsync.py
        └── animation/
            ├── __init__.py
            ├── facial_animation.py
            └── gesture_controller.py
    └── streaming/
        ├── __init__.py
        ├── kafka_producer.py
        ├── redis_producer.py
        └── event_schemas.py # Pydantic models for events
    └── utils/
        ├── __init__.py
        ├── audio_utils.py
        ├── video_utils.py
        └── logging_config.py
└── tests/
    └── test_perception.py
```

```
    └── test_avatar.py  
    └── test_webrtc.py  
  
    └── agent-core/          # LangGraph agents  
        ├── Dockerfile  
        ├── requirements.txt  
        └── src/  
            ├── __init__.py  
            ├── main.py          # FastAPI + LangGraph server  
            └── graphs/  
                ├── __init__.py  
                ├── orchestrator_graph.py  
                ├── dialogue_graph.py  
                └── reservation_graph.py  
            └── agents/  
                ├── __init__.py  
                ├── base_agent.py   # Abstract base class  
                └── orchestrator/  
                    ├── __init__.py  
                    ├── agent.py  
                    ├── session_router.py  
                    ├── backpressure_controller.py  
                    └── health_checker.py  
                └── dialogue/  
                    ├── __init__.py  
                    ├── agent.py  
                    ├── intent_classifier.py  
                    ├── context_manager.py  
                    ├── response_generator.py  
                    └── guardrails.py  
                └── reservation/  
                    ├── __init__.py  
                    ├── agent.py  
                    ├── table_matcher.py  
                    ├── hold_manager.py  
                    └── alternative_generator.py
```

```
    |   |   |   └── staff_notifier.py
    |   |   └── tools/
    |   |       ├── __init__.py
    |   |       ├── table_tools.py  # LangChain tools for table ops
    |   |       ├── knowledge_tools.py # RAG search tools
    |   |       ├── notification_tools.py
    |   |       └── calendar_tools.py
    |   └── memory/
        ├── __init__.py
        ├── session_memory.py
        ├── conversation_buffer.py
        └── long_term_memory.py
    └── rag/
        ├── __init__.py
        ├── embeddings.py
        ├── vector_store.py
        ├── retrievers.py
        └── knowledge_base.py
    └── streaming/
        ├── __init__.py
        ├── kafka_consumer.py
        ├── kafka_producer.py
        └── event_handlers.py
    └── utils/
        ├── __init__.py
        ├── llm_clients.py  # OpenAI, Claude, etc.
        ├── prompt_templates.py
        └── validation.py
└── tests/
    ├── test_orchestrator.py
    ├── test_dialogue.py
    └── test_reservation.py

└── backend-services/      # Supporting microservices
    └── table-management/
        └── Dockerfile
```

```
|- src/
|   |- main.py      # FastAPI CRUD service
|   |- models.py    # SQLAlchemy models
|   |- schemas.py   # Pydantic schemas
|   |- crud.py      # Database operations
|   |- database.py  # DB connection
|   |- api/
|       |- __init__.py
|       |- tables.py   # Table endpoints
|       |- reservations.py
|   |- migrations/   # Alembic migrations
|- knowledge-base/
|   |- Dockerfile
|   |- src/
|       |- main.py      # FastAPI RAG service
|       |- embeddings.py
|       |- vector_db.py
|       |- indexing.py
|       |- search.py
|   |- data/
|       |- menu.json
|       |- policies.md
|       |- faqs.json
|- notification-service/
|   |- Dockerfile
|   |- src/
|       |- main.py      # FastAPI notification service
|       |- sms_sender.py # SMS notifications
|       |- email_sender.py # Email notifications
|       |- slack_sender.py # Staff Slack alerts
|       |- webhook_sender.py
|   |- templates/
|       |- sms_templates.json
|       |- email_templates.html
|- web-frontend/      # React.js frontend
```

```
|   └── package.json  
|   └── tsconfig.json  
|   └── vite.config.ts  
|   └── public/  
|       └── index.html  
|       └── assets/  
|           └── avatars/      # 3D avatar models  
|               └── animations/    # Animation files  
|   └── src/  
|       └── main.tsx  
|       └── App.tsx  
|       └── components/  
|           └── Avatar/  
|               └── AvatarRenderer.tsx  
|               └── LipSyncController.tsx  
|               └── GestureController.tsx  
|           └── WebRTC/  
|               └── MediaCapture.tsx  
|               └── PeerConnection.tsx  
|               └── StreamManager.tsx  
|           └── UI/  
|               └── ChatInterface.tsx  
|               └── ReservationForm.tsx  
|               └── LoadingSpinner.tsx  
|           └── Admin/  
|               └── TableManagement.tsx  
|               └── ConversationLogs.tsx  
|               └── SystemHealth.tsx  
|       └── hooks/  
|           └── useWebRTC.ts  
|           └── useAvatar.ts  
|           └── useReservation.ts  
|       └── services/  
|           └── api.ts      # API client  
|           └── websocket.ts  # WebSocket connection  
|           └── webrtc.ts     # WebRTC utilities
```

```
    └── stores/      # State management
        ├── conversationStore.ts
        ├── reservationStore.ts
        └── avatarStore.ts
    └── types/
        ├── conversation.ts
        ├── reservation.ts
        └── webrtc.ts
    └── utils/
        ├── audioUtils.ts
        ├── videoUtils.ts
        └── formatters.ts
    └── dist/
    └── infrastructure/      # Deployment & DevOps
        └── kubernetes/
            ├── namespace.yaml
            └── edge-node/
                ├── deployment.yaml
                ├── service.yaml
                └── configmap.yaml
            └── agent-core/
                ├── deployment.yaml
                ├── service.yaml
                └── hpa.yaml      # Horizontal Pod Autoscaler
            └── backend-services/
                ├── table-management.yaml
                ├── knowledge-base.yaml
                └── notification-service.yaml
            └── databases/
                ├── postgres.yaml
                ├── redis.yaml
                └── kafka.yaml
            └── ingress/
                ├── nginx-ingress.yaml
                └── tls-certificates.yaml
```

```

└── terraform/
    ├── main.tf
    ├── variables.tf
    ├── outputs.tf
    ├── providers.tf
    └── modules/
        ├── gke-cluster/  # Google Kubernetes Engine
        ├── vpc-network/
        ├── load-balancer/
        └── monitoring/
    └── environments/
        ├── dev/
        ├── staging/
        └── production/
    └── monitoring/
        ├── prometheus/
        │   ├── config.yaml
        │   └── rules.yaml
        ├── grafana/
        │   ├── dashboards/
        │   └── datasources.yaml
        └── jaeger/          # Distributed tracing
            └── config.yaml
    └── scripts/
        ├── build.sh      # Build all services
        ├── deploy.sh     # Deploy to K8s
        ├── seed-data.sh  # Initialize databases
        └── health-check.sh  # System health verification

```

## 3. Event Schemas & API Contracts

### 3.1 Core Event Types (Pydantic Models)

```
# services/edge-node/src/streaming/event_schemas.py
from pydantic import BaseModel
```

```

from typing import Optional, List, Dict, Any
from datetime import datetime

class PerceptionEvent(BaseModel):
    type: str = "PerceptionEvent"
    session_id: str
    text: str
    start_ts: float
    end_ts: float
    speaker_id: Optional[str] = None
    emotion: Optional[str] = "neutral"
    asr_confidence: float
    language: str = "vi"# Vietnamese default
    metadata: Dict[str, Any] = {}

class AgentUtterance(BaseModel):
    type: str = "AgentUtterance"
    session_id: str
    text: str
    style: str = "friendly"
    ssml: Optional[str] = None
    gestures: List[str] = []
    emotion: str = "neutral"
    language: str = "vi"
    priority: int = 1# 1=normal, 2=urgent
    class ReservationRequest(BaseModel):
        session_id: str
        party_size: int
        preferred_time: datetime
        customer_name: Optional[str] = None
        contact_info: Optional[str] = None
        special_requests: List[str] = []

class SystemEvent(BaseModel):
    type: str
    session_id: str
    component: str

```

```
level: str# INFO, WARN, ERROR  
message: str  
timestamp: datetime  
metadata: Dict[str, Any] = {}
```

## 3.2 Database Schema

```
-- PostgreSQL schemasCREATE EXTENSION IF NOT EXISTS "uuid-ossp";  
CREATE EXTENSION IF NOT EXISTS "pgvector";  
  
-- Sessions tableCREATE TABLE sessions (  
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
    created_at TIMESTAMP DEFAULT NOW(),  
    updated_at TIMESTAMP DEFAULT NOW(),  
    status VARCHAR(20) DEFAULT 'active',  
    guest_count INTEGER DEFAULT 1,  
    language VARCHAR(5) DEFAULT 'vi',  
    metadata JSONB DEFAULT '{}'  
);  
  
-- Reservations tableCREATE TABLE reservations (  
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
    session_id UUID REFERENCES sessions(id),  
    table_id INTEGER,  
    party_size INTEGER NOT NULL,  
    reservation_time TIMESTAMP NOT NULL,  
    customer_name VARCHAR(100),  
    contact_info VARCHAR(100),  
    status VARCHAR(20) DEFAULT 'pending',  
    special_requests TEXT[],  
    created_at TIMESTAMP DEFAULT NOW(),  
    updated_at TIMESTAMP DEFAULT NOW()  
);  
  
-- Tables tableCREATE TABLE restaurant_tables (
```

```

        id SERIAL PRIMARY KEY,
        table_number VARCHAR(10) UNIQUE NOT NULL,
        capacity INTEGER NOT NULL,
        location VARCHAR(50),
        status VARCHAR(20) DEFAULT 'available',
        metadata JSONB DEFAULT '{}'
    );
-- Knowledge base for RAGCREATE TABLE knowledge_entries (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    title VARCHAR(200) NOT NULL,
    content TEXT NOT NULL,
    category VARCHAR(50),
    embedding vector(1536),-- OpenAI ada-002 dimensions
    metadata JSONB DEFAULT '{}',
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);
CREATE INDEX idx_knowledge_embedding ON knowledge_entries
USING ivfflat (embedding vector_cosine_ops);

```

## 4. Key Integration Points

### 4.1 WebRTC → Perception Pipeline

- **Protocol:** WebRTC DataChannel + MediaStream
- **Format:** Opus audio, VP9 video
- **Buffering:** Circular buffer for streaming ASR
- **Backpressure:** Drop frames if processing can't keep up

### 4.2 Perception → Agent Core

- **Transport:** Kafka topics or Redis Streams
- **Serialization:** JSON with Pydantic validation

- **Delivery:** At-least-once with deduplication
- **Batching:** Group events by session\_id for context

### 4.3 Agent Core → Avatar Pipeline

- **Transport:** Direct HTTP/2 calls or message queue
- **Caching:** Cache TTS audio for common phrases
- **Prioritization:** Queue with priority levels
- **Failover:** Fallback to simpler TTS if premium fails

### 4.4 Cross-Service Authentication

- **API Gateway:** Kong or Nginx with JWT validation
- **Service Mesh:** Istio for mTLS between services
- **Secrets:** Kubernetes secrets + HashiCorp Vault
- **Rate Limiting:** Per-session and per-service limits

## 5. Deployment Architecture

### 5.1 Edge Nodes (Low Latency)

- **Location:** Geographically close to restaurant
- **Services:** SFU + Perception + Avatar pipelines
- **Scaling:** Horizontal pod autoscaling based on WebRTC connections
- **Hardware:** GPU nodes for VAD/ASR inference

### 5.2 Core Nodes (High Compute)

- **Location:** Centralized cloud region
- **Services:** Agent Core + Backend Services
- **Scaling:** Based on message queue depth
- **Hardware:** CPU-optimized for LangGraph state machines

## 5.3 Data Layer

- **Primary DB:** PostgreSQL with read replicas
- **Cache:** Redis for session state + rate limiting
- **Message Bus:** Kafka cluster with replication
- **Object Storage:** S3/GCS for audio logs + models

This architecture provides a solid foundation for building a production-ready digital human restaurant assistant with clear separation of concerns, scalable components, and robust error handling.